

Lecture 7: Circuit Lower Bounds I

*Instructor: Rafael Pass**Scribe: J. Aaron Lenfestey*

We defined circuits in the last class and will spend the next two classes studying them. Circuits are an attractive area of study because they provide a level of structure which might be useful for proving that P and NP are separated. Recall the following definitions.

Definition 1 $\text{size}(T(n))$ is the class of languages deciable by circuit familes of size $O(T(n))$.

Definition 2 $P/\text{poly} = \bigcup_{c \in \mathbb{N}} \text{size}(n^c)$

There are two interesting alternative definitions of P/poly, as specified by the kinds of Turing machines that decide languages in the class:

Definition 3 A nonuniform polynomial time Turing machine is a sequence $\{M_n\}_{n \in \mathbb{N}}$ of Turing machines for which there exist a polynomial p such that $|M_n| \leq p(n)$ and M_n runs in time $p(n)$ on an input of length n .

Definition 4 A polynomial time Turing machine with advice is a Turing machine M and advice sequence $\{a_n\}_{n \in \mathbb{N}}$ for which there exists a polynomial p such that $M(x, a_n)$ runs in time $p(n)$ on inputs x of length n .

We are interested in P/poly because it is known that $P \subsetneq P/\text{poly}$ and believed that $P/\text{poly} \subsetneq NP$, and thus could potentially be used to find a separation between P and NP. We leave the former claim as an exercise and will provide evidence for the latter claim. Namely, if $NP \subseteq P/\text{poly}$, then the polynomial hierarchy collapses, which would indeed by a surprising turn of events.

Theorem 1 If $NP \subseteq P/\text{poly}$, then $PH = \Sigma_2$.

Proof. In particular, we will show that $NP \subseteq P/\text{poly} \implies \Pi_2 \subseteq \Sigma_2$, which we know implies $\Pi_2 = \Sigma_2$ and in turn that $PH = \Sigma_2$. Consider a polynomial time relation R and a Π_2 language L defined as:

$$x \in L \quad \text{iff} \quad \forall y_1 \exists y_2 . R(x, y_1, y_2).$$

We want to show that L is in Σ_2 . Since we are assuming $\text{NP} \subseteq \text{P/poly}$, we can construct a polynomial size circuit that *decides* SAT problems. Since finding SAT assignments can be done in polynomial time given a SAT oracle, there exists a polynomial q and $q(n)$ -sized circuit family $\{C_n\}_{n \in \mathbb{N}}$ that finds satisfying assignments to SAT problems. The trick is then to replace the inner existential quantification with the use of this circuit. Consider the following language L' :

$$x \in L' \quad \text{iff} \quad \exists C \, \forall y_1 \cdot |C| \leq q(x) \wedge R(x, y_1, C(x, y_1)).$$

Since the inner predicate checks that the circuit C is of polynomial size in x , it can be decided in polynomial time, so that the language L' is in Σ_2 . Now we show $L = L'$.

Suppose $x \in L$. Since R can be decided in time polynomial in x , it can be encoded into a polynomial size SAT problem (as in Carp's theorem). Then choose C to be a circuit as above that computes a satisfying assignment y_2 for $R(x, y_1, y_2)$ given x and y_1 . C is size polynomial in x , so that the predicate in the expression defining L' is satisfied.

Conversely, suppose $x \in L'$. The inner existential quantification in the definition of L would be satisfied by choosing $y_2 = C(x, y_1)$, so that $x \in L$.

Thus, we have shown that every language $L \in \Pi_2$ is in Σ_2 , so the polynomial hierarchy collapses. ■

This proof provides strong evidence that P/poly is a class that can distinguish P from NP . To do this, we would need to, for example, demonstrate that $\text{SAT} \notin \text{P/poly}$, which means that we are interested in establishing lower bounds on the size of circuits necessary for solving various problems.

Theorem 2 *Consider the class of functions $\{T_n^k\}$ on binary tuples x defined by $T_n^k(x_1, \dots, x_n) = 1$ iff $\sum x_i \geq k$. For $2 \leq k \leq n-1$, a boolean circuit of fan-in at most 2 that computes T_n^k must be of size at least $2n-4$.*

Proof. We will in all cases consider, without loss of generality, the smallest circuit that computes this function. This implies that there would be no circuit elements of fan-in 1 (e.g., unary negation) because one could create a smaller circuit by embedding this gate's functionality into downstream gates.

We proceed by induction on n . Consider the base case $n = 3 \implies k = 2$. Since every input must be examined to compute the function and each gate can examine only 2 inputs, there must be at least 2 gates in the network, which exactly meets the bound.

As our induction hypothesis, suppose that the theorem holds for

$$\{(n, k) \mid 3 \leq i < n \text{ and } 2 \leq k \leq i-1\}.$$

Consider a circuit C for computing T_n^k . Let G be a gate that takes two variables x_i and x_j as inputs for $i \neq j$ (recall that we have excluded unary gates). After fixing values for x_i and x_j , the function T_n^k restricted to the remaining $n-2$ variables will take on at least 3 different forms.

x_i	x_j	restricted form of T_n^k
0	0	T_{n-2}^k
0	1	T_{n-2}^{k-1}
1	0	T_{n-2}^{k-1}
1	1	T_{n-2}^{k-2}

Note that this is where the restriction $k < n$ becomes important. If $k = n$, then the restricted function takes only two forms, namely identically 0 or 1 if all its inputs are 1.

Since the restricted circuit assumes three different behaviors, however G only outputs a single boolean value, there must be a connection from either x_i or x_j to another gate G' in the circuit. Assume, without loss of generality, that the connection is from x_i . Now let C_α denote the circuit obtained by setting $x_i = \alpha$. Clearly, C_α computes $T_{n-1}^{k-\alpha}$. When constructing C_α we observe that G and G' are now unary gates, and so can be removed, so that $|C_\alpha| + 2 \leq |C|$. If $k = 2$, set $\alpha = 0$, other $\alpha = 1$ (so that $2 \leq k - \alpha \leq n - 2$). By the induction hypothesis, $|C_\alpha| \geq 2(n-1) - 4 \implies |C| \geq 2n - 4$. ■

Recall that a formula is a circuit where the fan-out of every node is 1. The circuit of a boolean formula is then just a binary tree.

Theorem 3 Consider the following function f_n defined on tuples \vec{x} of bit-strings of length $2\log(n)$:

$$f_n(x_1, \dots, x_n) = 1 \quad \text{iff} \quad \exists i \neq j \text{ s.t. } x_i = x_j.$$

Any formula computing f_n must be $\Omega(n^2)$.

Proof.

Since the circuits for formulas are trees, some input variables must be read in multiple times if their values must be processed by multiple gates. Let k be the total number of leaves (i.e., inputs) carrying bits from x_i , for some i , in a formula computing f_n . For some assignment $\vec{\alpha} = (\alpha_1, \dots, \alpha_n)$ to all the bits in \vec{x} , consider $f_n^{\vec{\alpha}-i}(x_i) = f_n(\alpha_1, \dots, x_i, \dots, \alpha_n)$.

After doing this restriction, the formula computing f_n can be reduced to a formula with k leaves that computes $f_n^{\vec{\alpha}-i}$ by eliminating unary gates. Thus the number of functions that $f_n^{\vec{\alpha}-i}$ can identically equal as α ranges over all possible values is bounded by the number of formulas of with k leaves. We count these.

$$\begin{aligned}
\# \text{ formulas on } k \text{ leaves} &\leq \# \text{ formulas of size } 2k & (1) \\
&\leq (\# \text{ binary trees of size } 2k) \cdot (\# \text{ ways to assign gates}) & (2) \\
&\leq 4^{2k} \cdot g^{2k} & (3) \\
&= 2^{O(k)} & (4) \\
& & (5)
\end{aligned}$$

Each x_i takes on $2^{2\log(n)} = n^2$ values. This means that, for various choices of α , the restricted function $f_n^{\bar{\alpha}-i}$ assumes at least $\binom{n^2}{n-1}$ functional forms (it is also sometimes identically 1). Therefore, it is necessary that

$$\binom{n^2}{n-1} \leq 2^{O(k)} \implies 2^n \leq 2^{O(k)} \implies k = \Omega(n).$$

Since the inputs in x_i need to be replicated $\Omega(n)$ times and this holds for each i , any circuit computing f_n must be of size $\Omega(n^2)$.